

11 - Naive Bayes

Text Classification

- **input** - a document d and a fixed set of classes $C = \{c_1, c_2, \dots, c_j\}$
- **output** - a predicted class $c \in C$
- *classification methods* - **hand-coded rules**
 - rules based on combinations of words or other features
 - accuracy can be high, if rules carefully refined by expert
 - but building and maintaining these rules is *expensive*
- **classification methods* - **supervised machine learning**
 - **input** - a document d , a fixed set of classes $C = \{c_1, c_2, \dots, c_j\}$, and a training set of m hand-labeled documents $(d_1, c_1), \dots, (d_m, c_m)$
 - **output** - a learned classifier $\gamma : d \Rightarrow c$
 - any kind of classifier
 - naive bayes
 - logistic regression
 - neural networks
 - k-nearest neighbors
 - etc

Naive Bayes Classifier

- **intuition for Naive Bayes** - simple classification method based on Bayes rule
 - relies on very simple representation of document, *bag of words*
- **Bayes' Rule** - for a document d and a class c , $P(c|d) = \frac{P(d|c)P(c)}{P(d)}$
- **Naive Bayes Classifier**
 - MAP = most likely class
 - $c_{MAP} = \arg \max (c \in C) P(c|d)$
 - using Bayes' Rule, can simplify to
 - $c_{MAP} = \arg \max (c \in C) P(d|c)P(c)$
 - where $P(d|c)$ is the *likelihood* and $P(c)$ is the *prior*
 - d can be represented as features x_1, \dots, x_n
 - $O(|X|^n \times |C|)$ parameters
 - could only be estimated if a very, very large number of training examples was available
 - we can just count the relative frequencies in a corpus

Multinomial Naive Bayes

- **independence assumptions** - $P(x_1, x_2, \dots, x_n | c)$
- **bag of words assumption** - assume position does not matter
- **conditional independence** - assume the feature probabilities $\{P(x_i | c_j)\}$ are independent given the class c
 - $P(x_1, x_2, \dots, x_n | c) = P(x_1 | c) \times P(x_2 | c) \times P(x_3 | c) \times \dots \times P(x_n | c)$
- $c_{NB} = \arg \max (c \in C) \prod (x \in X) P(x | c)$
- **problem** with multiplying lots of probabilities
 - can result in floating point underflow
 - *solution* - use logs since multiplications become additions
 - taking log does not change the ranking of classes
 - linear model - max of a sum of weights, so it is a linear function of the inputs
- Naive Bayes is a **linear classifier**

Learning

- *first attempt* - use MLE with the frequencies in the data
 - $P(w_i | c_j) = \frac{\text{count}(w_i, c_j)}{\sum_{(w \in V)} \text{count}(w, c_j)}$
 - fraction of times word w_i appears among all words in documents of topic c_j
 - create *mega-document* for topic j by concatenating all docs in this topic
 - use frequency of w in mega-document
 - **problem** - what if we have seen no training documents with a word classified in a class?
 - we will get 0
 - cannot condition away zero probabilities no matter what
 - **solution** - Laplace Add-1 smoothing, same idea as with Markov assumptions, add one to all counts
- **unknown words** - what do we do with them that appear in the test data but not in the training data or vocabulary?
 - **ignore them** - remove from test document, pretend they were not there
 - do not include any probability for them at all
 - building an unknown word model does not help, knowing which class has more unknown words is *not generally helpful*
- **stop words** - very frequent words like "the" and "a"
 - some systems ignore them
 - but usually does not help
 - sort vocabulary by word frequency in a training set
 - call the top 10 or 50 words the *stopword list*
 - remove all stop words from both training and test sets, as if they were never there to begin with

Relationship to Language Modeling

- **generative model of NB** - graph of the words that are classified to a specific class
- NB classifiers can use any sort of *feature*
 - i.e. URL, email address, dictionaries, network features
 - we use **only** word features
 - we use **all** of the words in the text, not a subset
 - then NB has an *important similarity* to language modeling
- **each class = a unigram language model**
 - assigning each word: $P(\text{word}|c)$
 - assigning each sentence: $P(s|c) = \prod P(\text{word}|c)$
 - example: each word and their probability it is positive class
 - I = 0.1, love = 0.1, this = 0.05, fun = 0.01, film = 0.1
 - $P(\text{sentence} | \text{positive}) = 0.1 \times 0.1 \times 0.05 \times 0.01 \times 0.1 = 0.0000005$
 - example: using positive class from previous and given new negative class, which one assigns the higher probability to sentence?
 - negative: I = 0.2, love = 0.001, this = 0.01, fun = 0.005, film = 0.1
 - $P(\text{sentence} | \text{negative}) = 10^{-9}$
 - $P(\text{sentence} | \text{positive}) > P(\text{sentence} | \text{negative})$

Naive Bayes Evaluation

- **2 by 2 confusion matrix**

		<i>gold standard labels</i>		
		gold positive	gold negative	
<i>system output labels</i>	system positive	true positive	false positive	precision = $\frac{tp}{tp+fp}$
	system negative	false negative	true negative	
		recall = $\frac{tp}{tp+fn}$		accuracy = $\frac{tp+tn}{tp+fp+tn+fn}$

- **accuracy** - do not use as the evaluation metric
 - useless, does not return what we are looking for
 - can get amazing accuracy for very dumb labeling that is not very representative of the data as a whole
 - use *precision* and *recall* instead
- **precision** - percent of items the system detected (i.e. items the system labeled as positive) that are positive (according to human gold labels)
 - $precision = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$
- **recall** - percent of items actually present in the input that were correctly identified by the system
 - $recall = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$

- precision and recall, not accuracy, emphasize **true positives** - finding the things that we are supposed to be looking for
- **combined measure F** - a single number that combines both *precision* and *recall*

$$F_{\beta} = \frac{(\beta^2 + 1)PR}{\beta^2 P + R}$$

-
- almost always use balanced F_1 ($\beta = 1$)
 - $F_1 = \frac{2PR}{P+R}$

Cross-Validation

- **devsets** - development test sets
- train on training set, tune on devset, report on test set
 - avoids *overfitting* - tuning to the test set
 - more conservative estimate of performance
 - **paradox** - want as much data as possible for training and as much for dev, so how do you split it?
- **cross-validation** - multiple splits
 - pool results over splits, compute pooled dev performance

Harms in Classification

- can have biases in classifiers that perpetuate negative stereotypes against a certain group of people, etc
 - or censorship of discussion about a group of people
- **causes of harms**
 - problems in the training data, ML systems known to amplify biases in their training data
 - problems in the human labels
 - problems in the resources used (like lexicons)
 - problems in the model architecture (like what the model is trained to optimize)
- **mitigation** of these harms is an *open research area*
- **model cards** - for each algorithm you release, document:
 - training algorithms and parameters
 - training data sources, motivation, and preprocessing
 - evaluation data sources, motivation, and preprocessing
 - intended use and users
 - model performance across different demographic or other groups and environmental situations